

Deep Q-Networks (DQN)

Your Name

Your Institution

July 19, 2025

Introduction to Deep Q-Networks (DQN)

Overview

Deep Q-Networks (DQN) are a groundbreaking approach in reinforcement learning that combines Q-learning with deep neural networks. This architecture allows agents to make decisions based on high-dimensional input data, like images.

- **Q-Learning:** Off-policy reinforcement learning algorithm that estimates the action-value function $Q(s, a)$ to learn the value of action choices in various states.
- **Deep Learning:** Utilizes neural networks with multiple layers to extract features automatically; in DQN, it approximates the Q-function, learning policies from visual inputs.

Architecture of DQN

- 1 **Input Layer:** Takes pre-processed state representations (e.g., frames from a game).
- 2 **Hidden Layers:** One or more fully connected layers with nonlinear activation functions to capture complex features.
- 3 **Output Layer:** Outputs Q-values for each possible action, selecting the one with the highest Q-value.

Significance of DQN

- **Handling High Dimensionality:** Successfully processes visual input, applicable for games.
- **Experience Replay:** Stores past experiences in a memory buffer, allowing efficient learning.
- **Target Network:** A separate network providing a stable target Q-value for main network updates.

Example of DQN in Action

Example

An agent plays a video game by feeding frames into the DQN, processing visual information to predict Q-values for actions (e.g., move left, jump, shoot), and updating knowledge based on rewards from actions taken.

Key Points to Emphasize

- DQN is an advancement over traditional Q-learning utilizing deep learning for enhanced performance.
- The combination of experience replay and target networks is vital for stable learning.
- DQN has demonstrated success in various domains, particularly in Atari games.

Q-Learning Update Rule

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right] \quad (1)$$

- Where α is the learning rate, r is the reward, γ is the discount factor, and s' is the next state.

Conclusion

DQN marks a pivotal moment in reinforcement learning, effectively using deep neural networks to tackle complex decision-making tasks and advancing AI in high-dimensional observation environments.

Fundamentals of Q-Learning - Overview

- Q-Learning is a model-free reinforcement learning algorithm.
- It helps agents learn to select optimal actions based on their experiences.
- This slide reviews Q-Learning, its functions, limitations, and how DQNs propose to overcome these challenges.

What is Q-Learning?

- Q-Learning is a model-free reinforcement learning algorithm.
- It learns the value of actions in states, enabling decision-making.

Q-Value (Action-Value Function)

Represents the expected future rewards of taking action a in state s , denoted as $Q(s, a)$.

- **Learning Rate (α):** Value between 0 and 1 that determines how much of the new Q-value overwrites the old one.
- **Discount Factor (γ):** Value between 0 and 1 that balances long-term and short-term rewards.

Q-Learning Update Rule

The Q-value is updated using the formula:

$$Q(s, a) \leftarrow Q(s, a) + \alpha \left(r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right) \quad (2)$$

- s' : next state after taking action a .
- r : reward received after state transition from s to s' .

Example of Q-Learning

- Imagine an agent navigating a simple maze:
 - **States (s):** Each position in the maze.
 - **Actions (a):** Move Up, Down, Left, Right.
- The agent explores different actions, updates its Q-values based on rewards, learning the best path to the goal.

Limitations of Q-Learning

- **Scalability:** Q-table grows exponentially with the number of states and actions, making it impractical in large state spaces.
- **Exploration vs. Exploitation:** Finding the right balance can be challenging.
- **Convergence:** Requires significant training data to reach the optimal policy, which can be slow.

How DQN Addresses Limitations

- Uses deep neural networks to approximate Q-values instead of storing them, handling large state spaces.
- Introduces **experience replay** to reuse past experiences, improving learning efficiency.
- Utilizes **target networks** for stable training, reducing issues from correlated updates.

Introduction to Deep Q-Networks (DQN)

- Definition: Combines Q-learning with deep learning techniques.
- Purpose: Enables effective approximation of Q-value functions using deep neural networks.

Key Concepts in DQN

- **Q-Learning:**
 - Model-free reinforcement learning algorithm.
 - Learns value of actions in given states.
- **Deep Learning:**
 - Utilizes deep neural networks to learn patterns from complex data.
- **Function Approximation:**
 - DQNs use deep networks to estimate Q-values, solving issues with vast state spaces.

Example Illustration and Key Points

- **Example:** A robot navigating a maze
 - Traditional Q-learning: Uses a fixed Q-table for specific positions.
 - DQN: Employs a neural network for Q-value approximation from maze images.
- **Key Points:**
 - 1 Bridges Q-learning and deep learning for efficient complex data learning.
 - 2 Simplifies high-dimensional problems by approximating Q-values.
 - 3 Uses **Experience Replay** for stabilizing training.

Summary and Next Steps

- DQNs represent a significant advancement in reinforcement learning.
- They blend Q-learning's strengths with deep learning's capacity for complex inputs.
- **Next Step:** Explore the architecture of DQNs.

Q-Learning Update Rule

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left(r_t + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t) \right) \quad (3)$$

Where:

- s_t : Current state
- a_t : Action taken
- r_t : Reward received
- α : Learning rate
- γ : Discount factor

Overview of DQN

Deep Q-Networks (DQN) merge traditional Q-learning with deep learning, enabling AI to learn effective policies from high-dimensional input spaces (like images). Understanding the architecture is essential for grasping how DQNs operate.

Architecture of DQN - Input and Hidden Layers

① Input Layer:

- **What it Accepts:** Current state representation of the environment (e.g., raw pixels in a game).
- **Data Preprocessing:** Transformations (e.g., grayscale conversion, resizing) reduce computational overhead.

② Hidden Layers:

- **Convolutional Layers:**
 - **Purpose:** Detect features from the input state (e.g., edges, shapes).
 - **Functionality:** Filters generate feature maps highlighting important spatial hierarchies.
- **Fully Connected Layers:**
 - **Purpose:** Combine features learned by convolutional layers.
 - **Activation Functions:** Typically, Rectified Linear Unit (ReLU) for non-linear transformations.

Architecture of DQN - Output Layer and Key Points

- **Output Layer:**

- **What it Produces:** Q-values for each possible action in the current state.
- **Size of Output:** If there are 4 actions, the output layer has 4 nodes.
- **Final Output:** The Q-value estimates expected future rewards for each action.

Key Points to Emphasize

- Efficient learning from high-dimensional inputs.
- Convolutional layers are critical for feature extraction.
- Experience replay enhances learning efficiency.

Architecture of DQN - Code Snippet Example

```
class DQNModel(nn.Module):
    def __init__(self, action_size):
        super(DQNModel, self).__init__()
        self.conv1 = nn.Conv2d(in_channels=1,
                                out_channels=32, kernel_size=8, stride=4)
        self.conv2 = nn.Conv2d(in_channels=32,
                                out_channels=64, kernel_size=4, stride=2)
        self.fc1 = nn.Linear(64 * 7 * 7, 1024)
        self.output = nn.Linear(1024, action_size)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        x = F.relu(self.conv2(x))
        x = x.view(x.size(0), -1)    # Flatten the
                                     output
        x = F.relu(self.fc1(x))
        return self.output(x)
```

Definition

Experience Replay is a technique used in Reinforcement Learning (RL), particularly in Deep Q-Networks (DQN), where an agent stores past experiences to train the model. By reusing previous experiences, the agent learns more efficiently.

Experience Replay - How It Works

1 Storage:

- The agent collects experiences as tuples: (s_t, a_t, r_t, s_{t+1})
- Where:
 - s_t = state at time t
 - a_t = action taken at state s_t
 - r_t = reward received for action a_t
 - s_{t+1} = new state after taking action a_t

2 Replay Buffer:

- Experiences are stored in a ****replay buffer**** with a fixed size.
- Older experiences are discarded when new experiences fill it up.

3 Sample & Train:

- Batches of experiences are randomly sampled to update the DQN.
- This helps break the correlation between consecutive experiences and stabilizes training.

Experience Replay - Advantages and Example

Advantages of Experience Replay

1 Efficiency in Learning:

- Reusing experiences allows the agent to learn from past actions multiple times, accelerating the learning process.

2 Increased Stability:

- Breaking correlation reduces update variance, leading to more stable training.

3 Diverse Training:

- Sampling from varied experiences helps the model generalize better.

Illustrative Example

Consider an agent playing a game:

- Experience: (s_t, a_t, r_t, s_{t+1}) :
 - s_t = "Player at position (3, 4)"
 - a_t = "Move Right"
 - r_t = "+10 points"

Target Network Mechanism in DQN - Overview

Overview of the Target Network

In Deep Q-Networks (DQN), the target network stabilizes training by reducing correlations among Q-value updates. This mechanism mitigates oscillations and divergence that can occur when learning from rapidly changing targets.

How the Target Network Works

① Architecture:

- DQN uses two neural networks: the online and the target network, having the same architecture but updated at different intervals.

② Q-value Updates:

- The online network generates Q-values based on the current state while the target network provides stable targets.

③ Updating the Target Network:

- The target network's weights are updated to match the online network every fixed number of steps (e.g., every 1000 steps), ensuring stable targets.

Target Network Mechanism in DQN - Impact and Example

Key Points to Emphasize

- **Stability:** Provides a consistent target over multiple updates.
- **Delay in Updates:** Acts as a buffer against sudden Q-value changes.
- **Effect on Learning:** Separate networks enhance convergence and performance.

Example Scenario

Suppose the online network predicts Q-values for actions A, B, and C. The target Q-values from the target network guide the updates of the online network.

Code Snippet

```
# Pseudo code for updating target network
if step % TARGET_UPDATE_FREQ == 0:
    target_network.load_state_dict(online_network.
        state_dict())
```

Target Network Mechanism in DQN - Summary

Summary

The target network in DQN significantly enhances training stability by decoupling the learning process of Q-values from their targets. This leads to more reliable updates and improved performance in reinforcement learning tasks.

Loss Function in DQN - Overview

In Deep Q-Networks (DQN), the loss function is crucial for updating Q-values and improving decision-making. Understanding this loss function is essential for grasping how DQN optimizes its learning process.

Key Concepts - Q-Learning and Loss Function

1 Q-Learning

- A reinforcement learning algorithm to learn the value of actions in states.
- Goal: Approximate the optimal action-value function, $Q^*(s, a)$.

2 Loss Function

- Used to measure the difference between predicted and target Q-values.
- Defined as Mean Squared Error (MSE):

$$L(\theta) = \mathbb{E}_{(s,a,r,s')} \left[(r + \gamma \max_{a'} Q(s', a'; \theta^-) - Q(s, a; \theta))^2 \right] \quad (4)$$

- Where:
 - r : Reward after taking action a in state s
 - γ : Discount factor (importance of future rewards)
 - $Q(s, a; \theta)$: Predicted Q-value
 - $Q(s', a'; \theta^-)$: Target Q-value from target network

Key Concepts - Target Network and How DQN Minimizes Loss

3 Target Network

- A secondary neural network updated less frequently for stability.

4 How DQN Minimizes Loss

- **Experience Replay**
 - Stores previous experiences to sample during training for stability.
- **Stochastic Gradient Descent (SGD)**

$$\theta \leftarrow \theta - \alpha \nabla L(\theta) \quad (5)$$

- α : Learning rate
- **Convergence**
 - Adjusts Q-values to minimize loss for optimal policy decisions.

Example and Key Points

Example: Suppose an agent receives a reward of +1 after choosing action a in state s and transitioning to state s' .

- Predicted $Q(s, a; \theta) = 0.5$ and target $Q(s', a'; \theta^-) = 0.8$.
- Target calculation:

$$target = r + \gamma \max_{a'} Q(s', a'; \theta^-) = 1 + 0.9 \times 0.8 = 1.72 \quad (6)$$

- Loss calculation:

$$L(\theta) = (1.72 - 0.5)^2 = 1.2976 \quad (7)$$

Key Points:

- The loss function is central to the learning process in DQN.
- Target networks and experience replay are vital for stable learning.

Training Process - Overview

In this slide, we will explore the steps involved in training a Deep Q-Network (DQN), focusing on:

- Epochs
- Batch Updates
- Convergence

1. Epochs

- **Definition:** An epoch refers to one complete pass through the entire training dataset. In the context of DQN, it involves using episodes of interaction with the environment.
- **Example:** For example, if training a DQN to play a video game, an epoch could involve the agent playing through 100 games, experiencing different states, taking actions, and receiving rewards.

2. Batch Updates

- **Experience Replay:** This technique helps break the correlation between consecutive experiences by storing experiences in a replay buffer.
- **Mini-Batch Size:** Typically, a mini-batch size of 32 or 64 is used to stabilize training and improve convergence.
- **Loss Function Calculation:**

$$L(\theta) = \mathbb{E}_{(s,a,r,s') \sim D} \left[(y - Q(s, a; \theta))^2 \right] \quad (8)$$

where:

- $y = r + \gamma \max_{a'} Q(s', a'; \theta^-)$
- D : replay buffer
- θ : parameters of the current Q-network
- θ^- : parameters of the target network

3. Convergence

- **Goal:** The goal of training is to converge the Q-values to the optimal Q-values that reflect the expected returns of actions taken in given states.
- **Monitoring Convergence:** This is gauged by observing if the loss stabilizes and if the Q-values approach a fixed point over time.
- A common practice is to plot the average loss over several epochs to visualize convergence.

Training Process - Summary

The training process of DQN includes:

- **Epochs:** Interaction with the environment
- **Batch Updates:** Utilizing experience replay for stable learning
- **Convergence:** Indicating successful learning

Understanding these components is essential for properly training DQNs and applying them effectively in various scenarios.

Applications of Deep Q-Networks (DQN)

Introduction to DQN Applications

Deep Q-Networks (DQN) integrate deep learning into reinforcement learning, effectively handling high-dimensional state spaces and learning optimal policies from raw sensory data.

- **Industry Impact:** DQN outperformed human players in complex video games.
- **Example: Atari Games**
 - *Game: Breakout*
 - DQN learned to play and exceeded human-level performance from pixel data.
 - Optimized strategies for ball trajectory and paddle positioning using experience replay.
- **Key Highlights**
 - Achieved superhuman performance in several games.
 - Utilizes frame stacking to capture temporal dynamics.

Applications of DQN - Robotics, Autonomous Vehicles, and Finance

• Robotics

- DQNs control robots in complex environments.
- *Example: Robotic Manipulation*
 - Task: Control a robotic arm to pick up objects.
 - Adapts well to new configurations and object types.

• Autonomous Vehicles

- Application in self-driving cars for decision-making.
- *Example: Path Planning*
 - Helps in lane changing and obstacle avoidance using past experiences.

• Finance

- Aids in algorithmic trading strategies.
- *Example: Stock Trading*
 - Makes buy/sell decisions based on historical market conditions.

Conclusion and Key Takeaway

- DQNs show versatility across gaming, robotics, autonomous vehicles, and finance.
- They effectively tackle complex decision-making tasks.
- **Key Takeaway:** DQNs represent a leap in artificial intelligence, showcasing the power of combining reinforcement learning with deep learning.

Challenges Faced by DQN - Part 1

1 Overestimation Bias:

- Description: DQNs often tend to overestimate action values due to function approximation.
- Illustration: Q-values for certain actions get exaggerated, leading to suboptimal policy decisions.
- Example: In a game, a high Q-value for an unfavorable action may cause the agent to prefer it over better options.

2 Instability and Divergence:

- Description: Training can become unstable from correlations in data and continuous updates to the Q-network.
- Example: Small perturbations in input may lead to drastic output changes if the model is not converged.

③ Sample Inefficiency:

- Description: DQNs require significant experience for training due to high dimensional state spaces.
- Example: In complex environments, learning optimal policies may take thousands of episodes.

④ Need for Hyperparameter Tuning:

- Description: Multiple hyperparameters (learning rate, discount factor) heavily influence training success.
- Key Point: Finding optimal hyperparameters necessitates extensive trial and error, which is time-consuming.

Future Research Opportunities in DQN

① Improved Value Function Estimation:

- Goal: Reduce overestimation bias using techniques like Double DQN.
- Future Directions: Explore ensemble approaches for multiple value estimators to improve accuracy.

② Algorithmic Enhancements:

- Ideas: Investigate methods like Dueling DQN for better training efficiency using separate state value and advantage estimates.
- Focus: Enhancing DQN architecture for more robust policies with lesser data.

③ Transfer Learning and Meta-Reinforcement Learning:

- Description: Integrate knowledge from previous tasks to expedite new task training.
- Future Directions: Apply DQNs in dynamic real-world settings, learning from past experiences.

④ Multi-Agent Systems:

- Description: Extend DQNs for cooperation among multiple agents learning concurrently.
- Future Directions: Explore communication strategies and collaboration to improve learning in complex environments.